

Mobile Phones Price: a classification problem

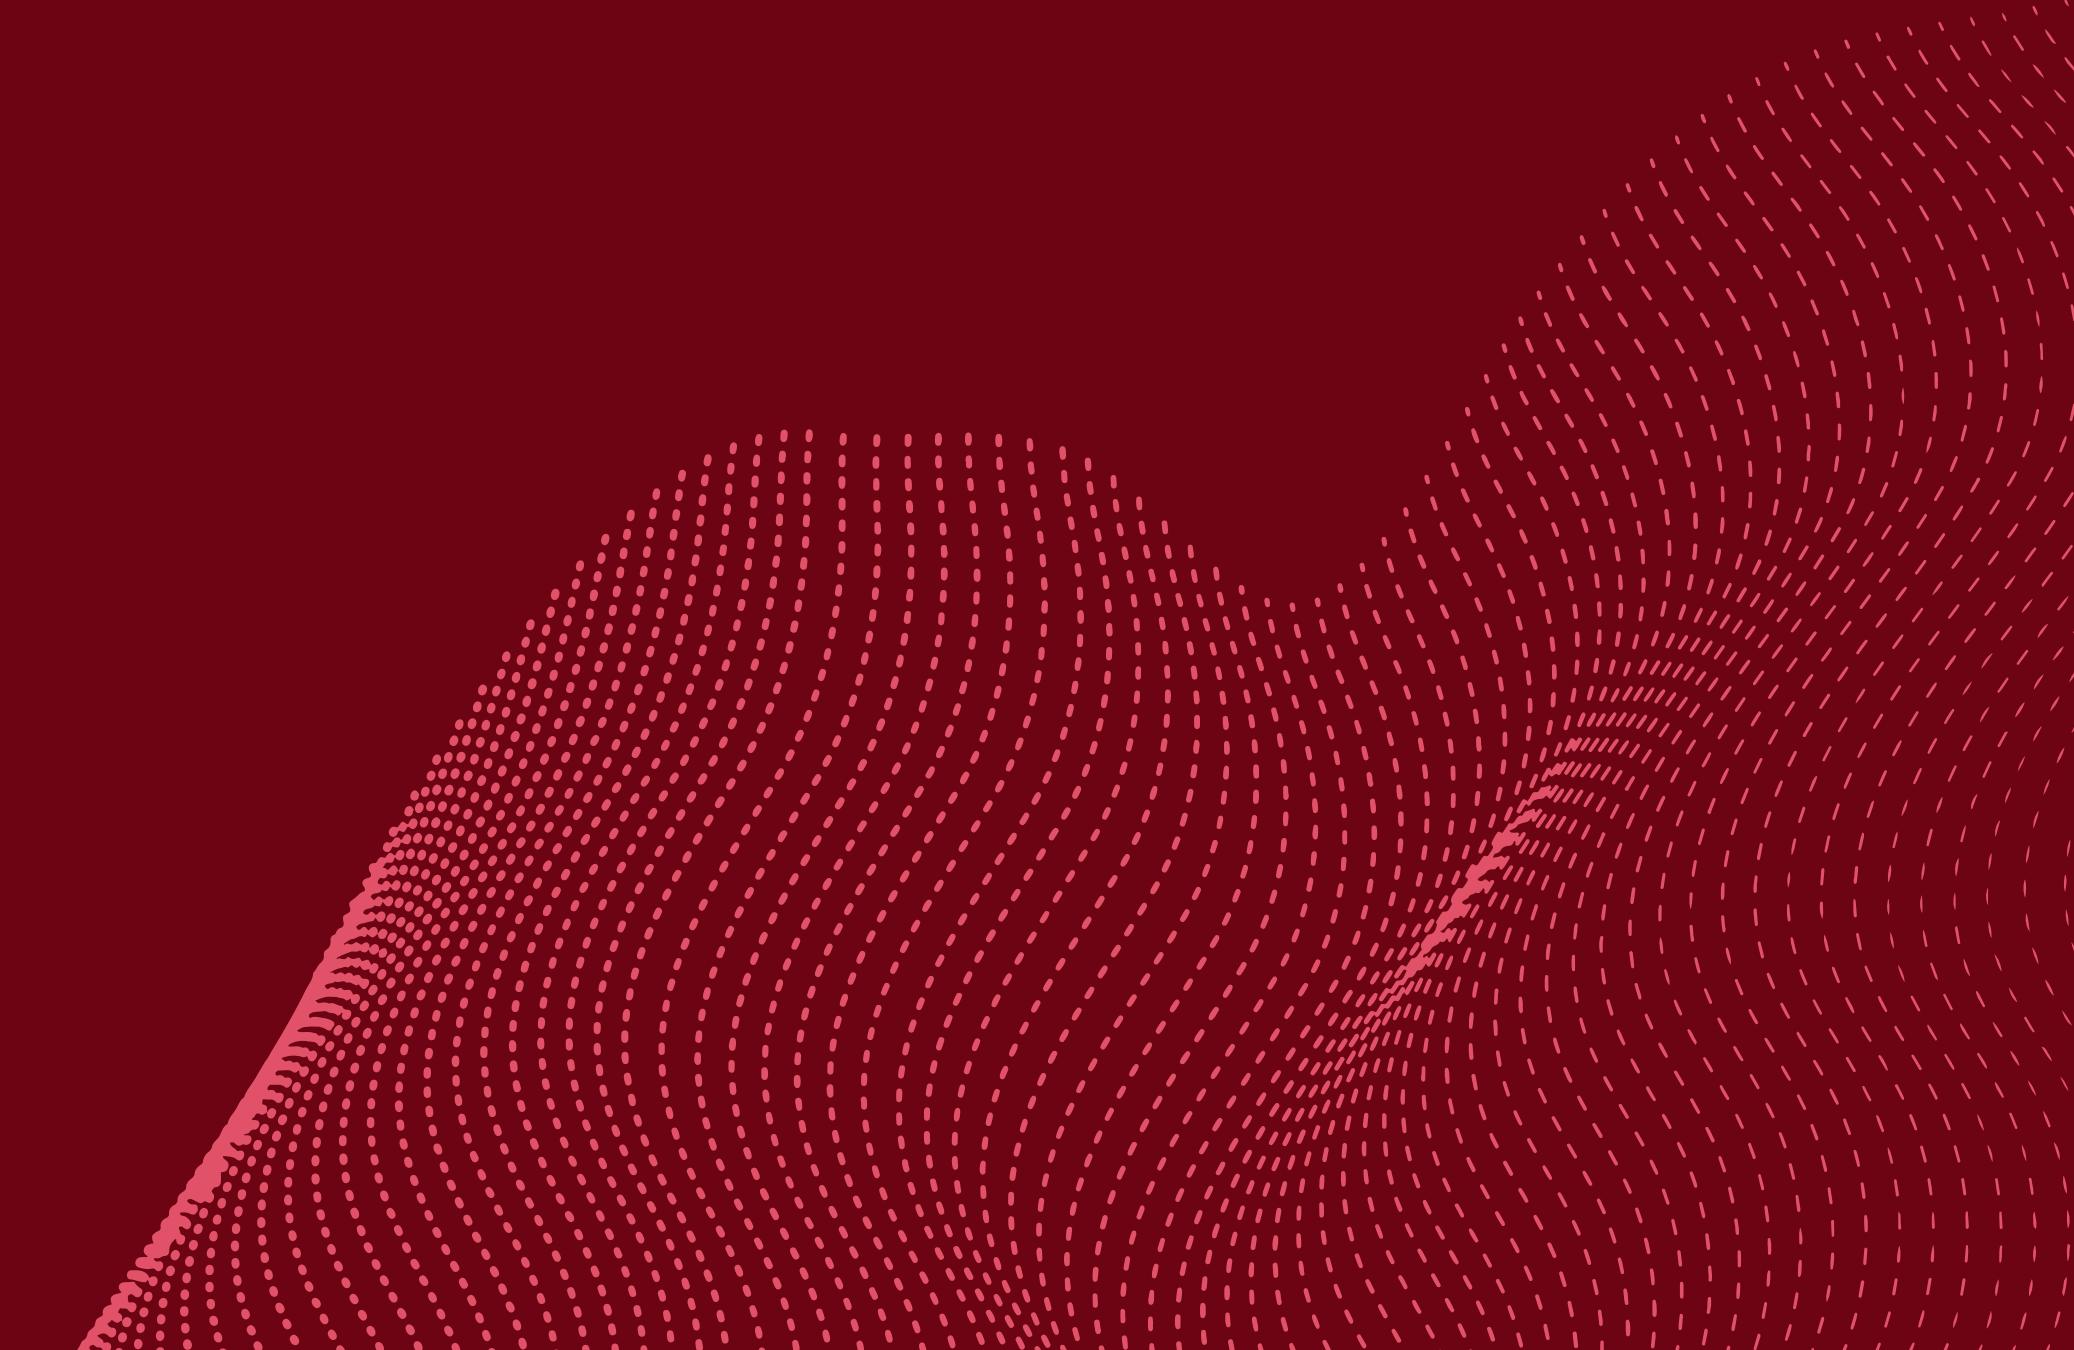
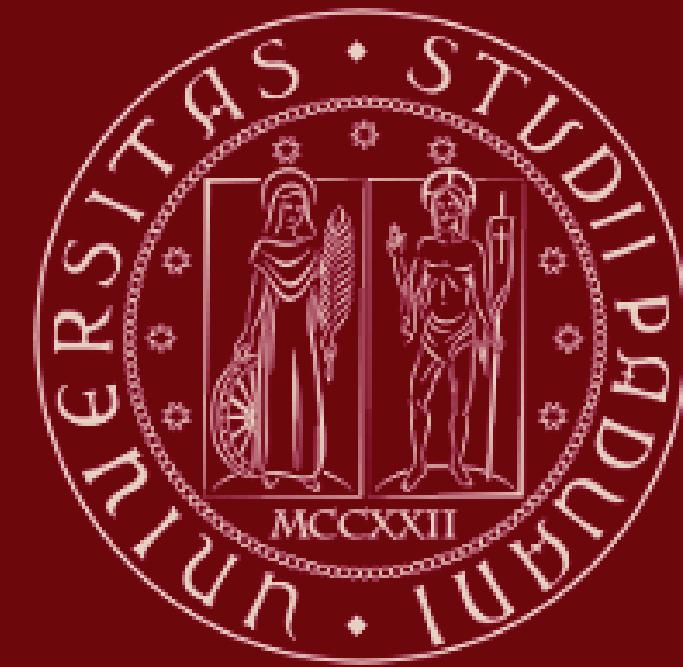
Statistical Learning course

Anthony Palmieri

Pietro Maria Sanguin

Angelica Giangiacomi

Academic year: 2021/2022



Overview

- Dataset: Mobile Price Train
- Preprocessing and Data Cleaning
- Exploratory Data Analysis
- Training Phase
 - ↳ Multinomial Logistic Regression
 - ↳ Linear Discriminant Analysis
 - ↳ Quadratic Discriminant Analysis
 - ↳ K-Nearest Neighbors
- Conclusions



Dataset: Mobile Price Train

→ Multiclass Classification Problem

↳ 4 labels: low-cost(0), medium-cost(1), high-cost(2), very-high cost(3)

→ 2000 examples

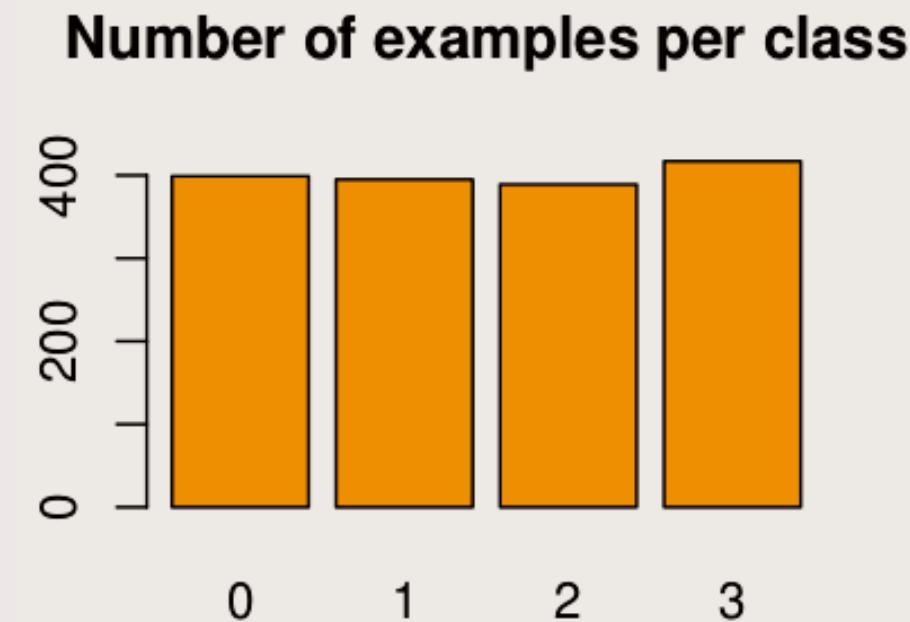
→ 20 features

↳ Categorical: blue, dual_sim, four_g, three_g, touch_screen, wifi

↳ Numerical: battery_power, clock_speed, fc, int_memory, m_depth,
mobile_wt, n_cores, pc, px_height, px_width, ram, sc_h,
sc_w, talk_time

Preprocessing and Data Cleaning

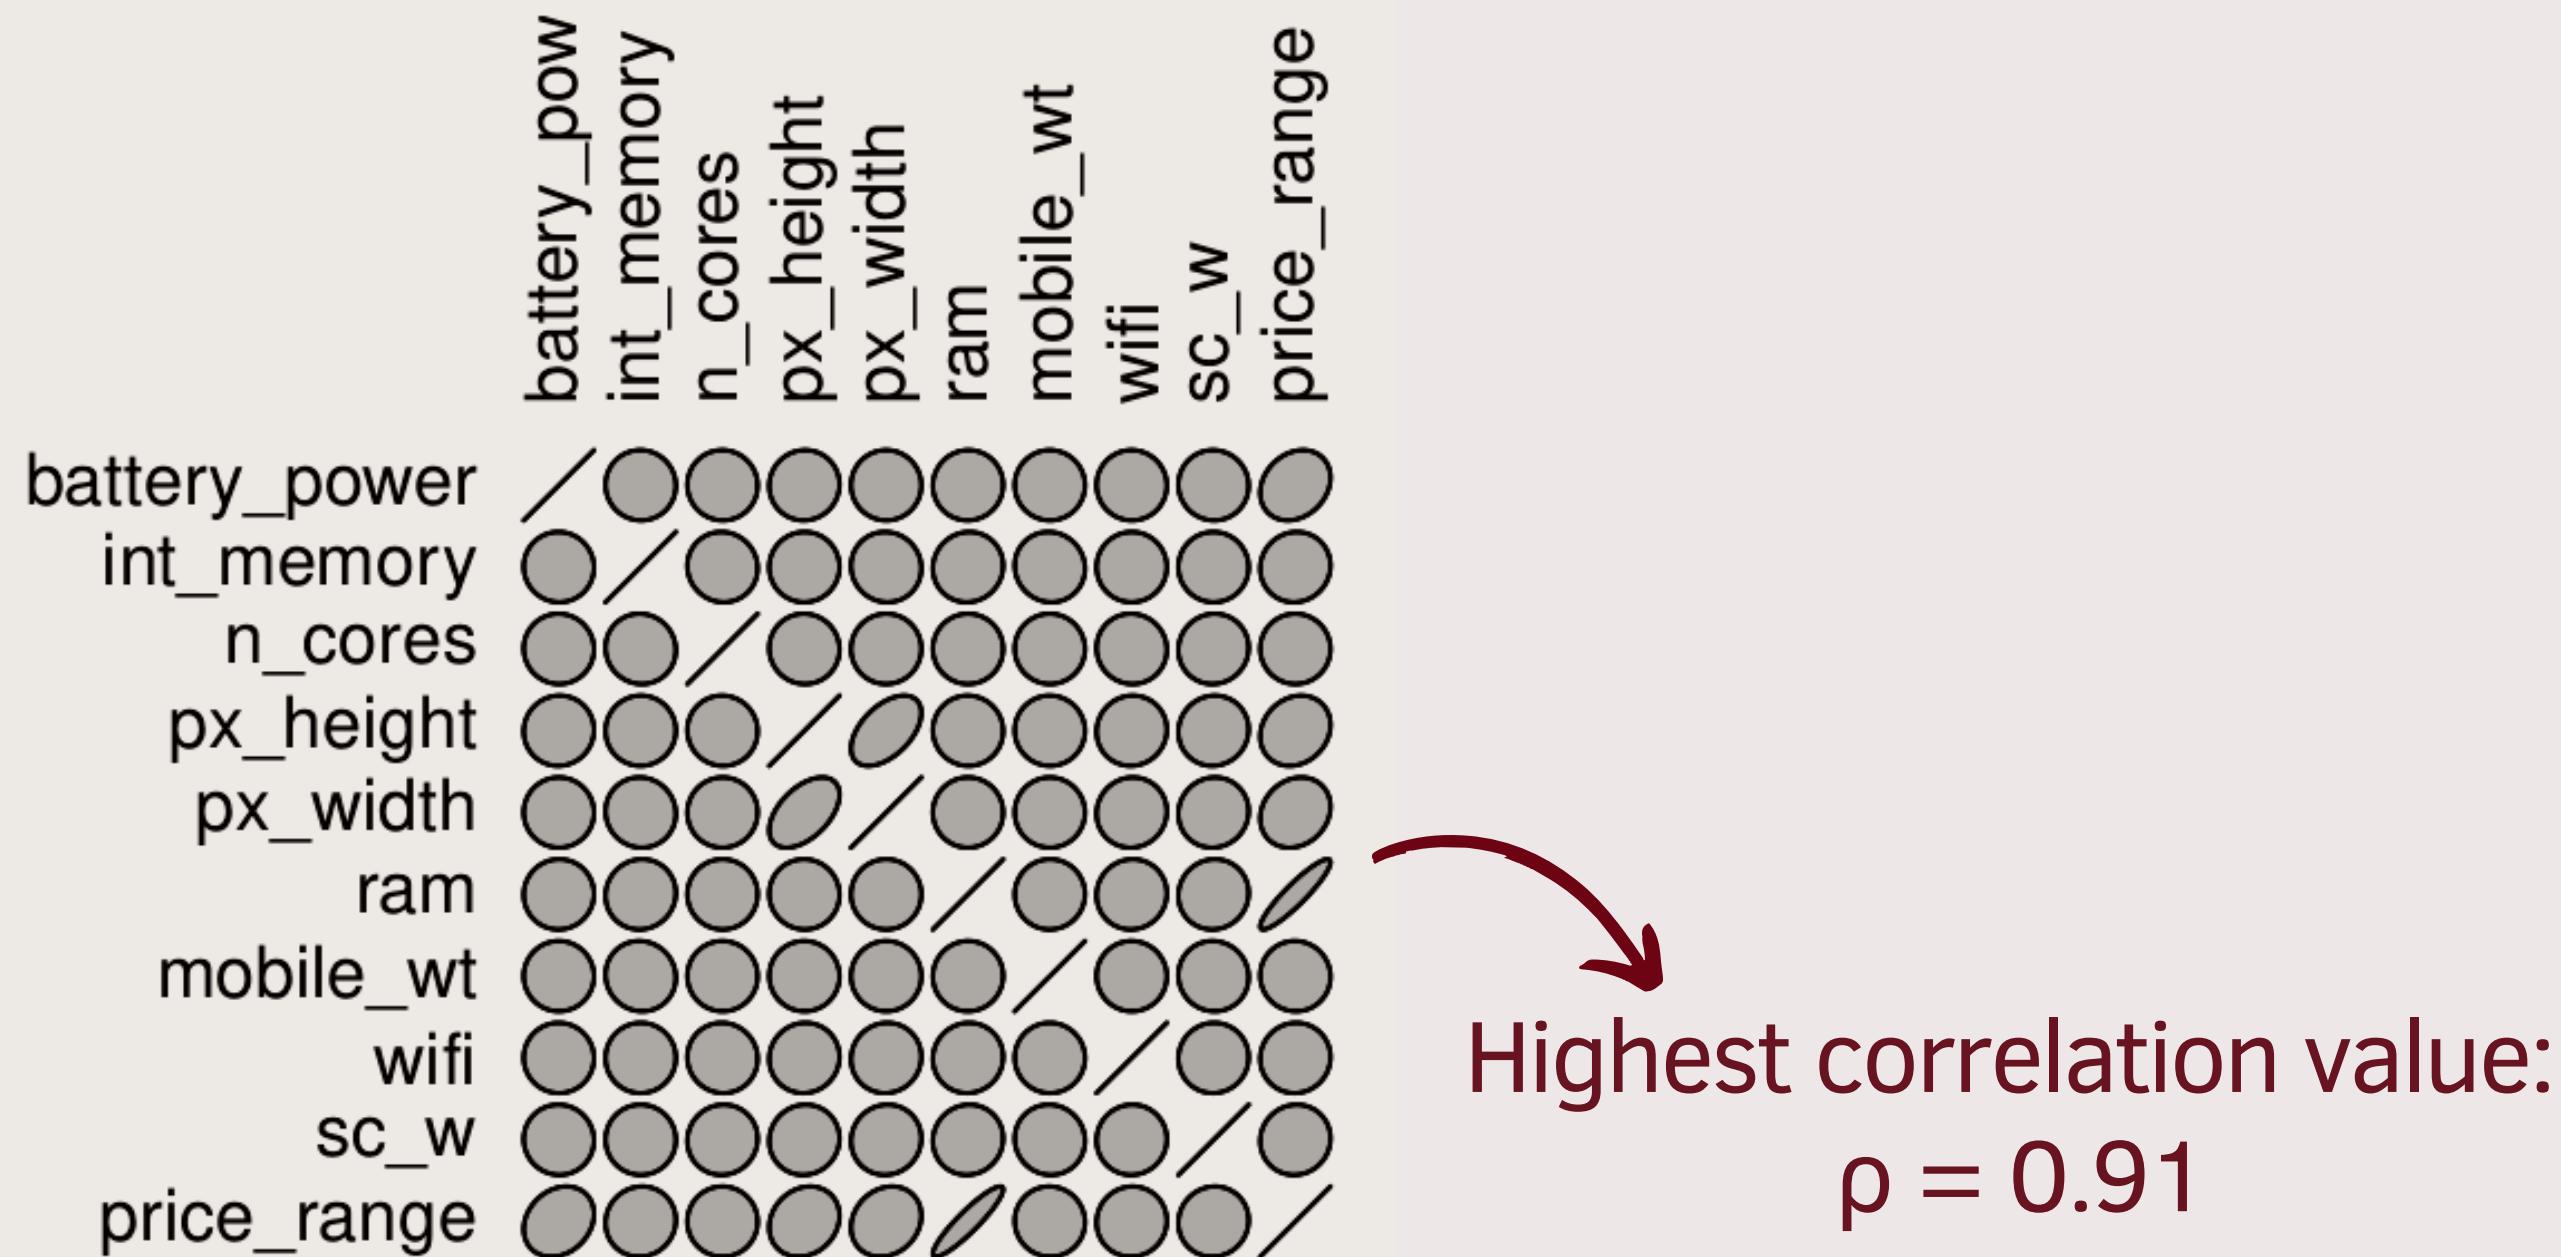
- Not Available or Not a Number values? X
- Split dataset into Training and Test sets
- Check for Class Balance



- Feature Scaling: MinMax Scaler

Exploratory Data Analysis (I)

→ Correlation Ellipses



Exploratory Data Analysis (II)



Exploratory Data Analysis (III)

→ Check for Multicollinearity using Variance Inflation Factor

```
# print all VIFs values  
vifs
```

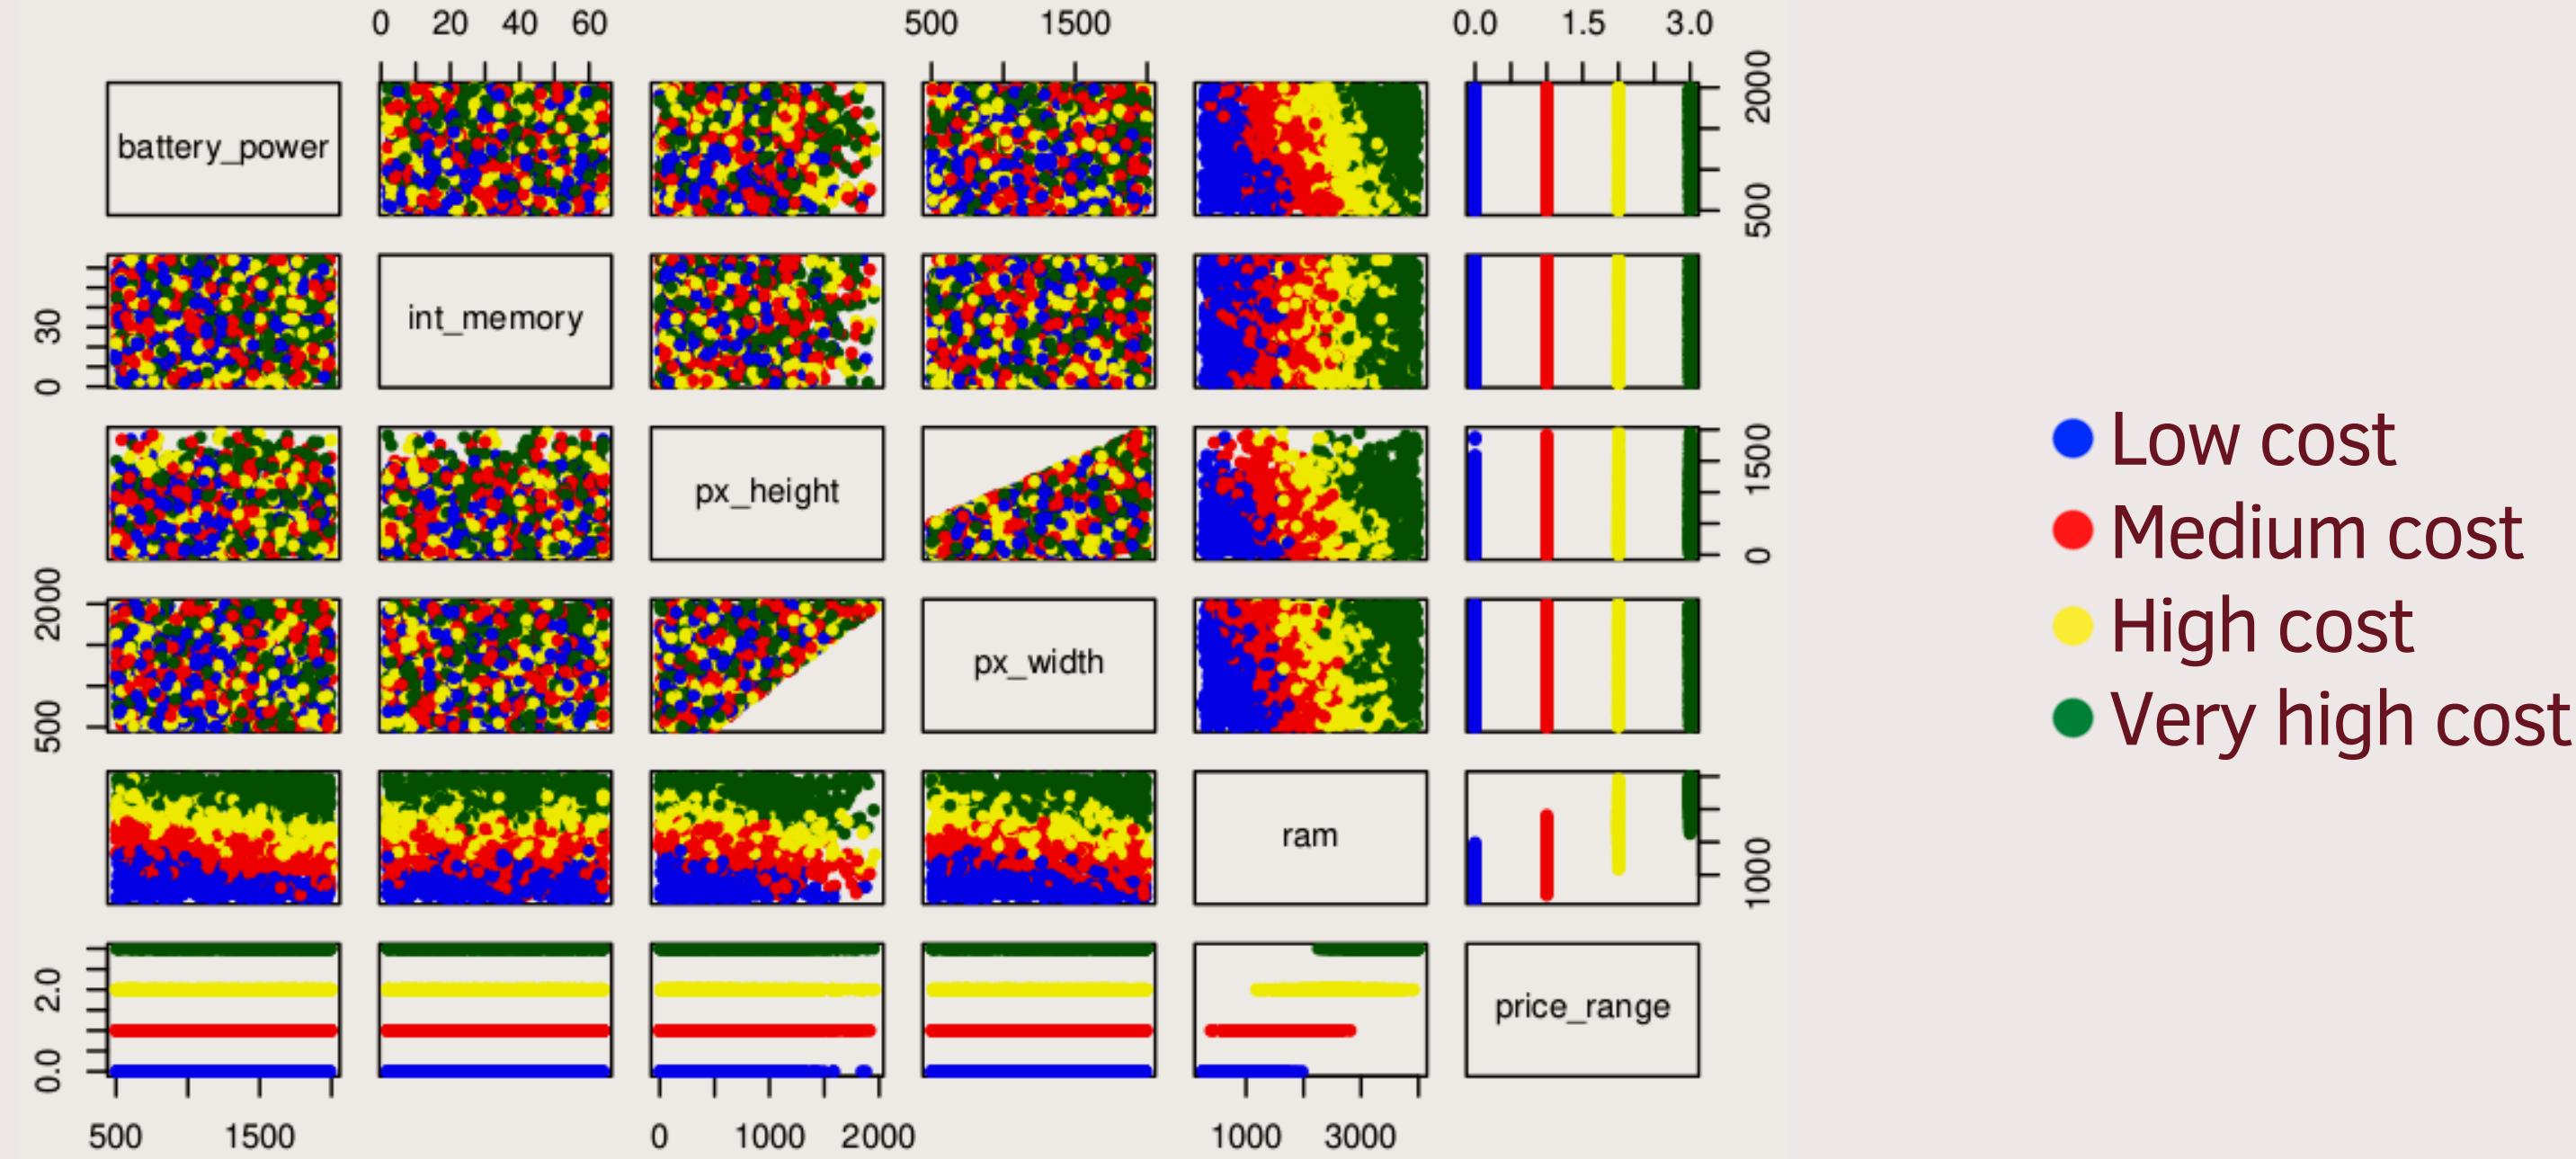
```
## [1] 1.011139 1.002891 1.730141 1.010976 1.007034 1.003561 1.006990 1.736878  
## [9] 1.353688 1.352578 1.006079 1.399044 1.396769 1.006836
```



All values close to 1 ↽ no important multicollinear relationships

Exploratory Data Analysis (IV)

→ Check for Non-Linear Relationships



no important Non-Linear relationships

Training Phase

→ Multinomial Logistic Regression

~~~ a classification method that generalizes Logistic Regression to multiclass problems

## → Linear Discriminant Analysis

~~~ a classification method that under normality assumptions uses Bayes theorem to derive discrimination scores

→ Quadratic Discriminant Analysis

~~~ a classification method just like LDA but does not assume equal variance among the various groups

## → K - Nearest Neighbors

~~~ a classification method that uses a majority scheme on the K nearest examples to classify each point

Multinomial Logistic Regression (Full Dataset)

→ Train the model keeping all features

```
# original model with all predictors  
  
# create formula to fit the model  
features = paste(names(X_train), collapse='+')  
formula = paste(c('y_train', features), collapse='~')  
  
# Fit the Multinomial Logistic Regression model  
mlr_full <- nnet::multinom(formula, data = X_train_scaled)
```



```
# Make predictions on Training set  
y_train_pred <- predict(mlr_full, X_train_scaled)  
  
# Model accuracy on Training set  
cat("Training set accuracy: ", mean(y_train_pred == y_train), "\n")  
  
## Training set accuracy: 0.995  
  
# Make predictions on Test set  
y_test_pred <- predict(mlr_full, X_test_scaled)  
  
# Model accuracy on Test set  
cat("Test set accuracy: ", mean(y_test_pred == y_test))  
  
## Test set accuracy: 0.975
```

Multinomial Logistic Regression (BWS)

→ Train the model using p-values to apply Backward Selection

```
# Final model after BACKWARD SELECTION

# create formula to fit the model
features = paste(names(X_train)[-c(19, 18, 2, 8, 6, 15, 17, 4, 16, 5, 3, 11, 20)], collapse='+')
formula = paste(c('y_train', features), collapse='~')

# Fit the Multinomial Logistic Regression model
mlr_BS <- nnet::multinom(formula, data = X_train_scaled)
```

Features kept:
battery_power, int_memory, mobile_wt,
n_cores, px_height, px_width, ram



```
# Make predictions on Training set
y_train_pred <- predict(mlr_BS, X_train_scaled)

# Model accuracy on Training set
cat("Training set accuracy: ", mean(y_train_pred == y_train), "\n")

## Training set accuracy: 0.985

# Make predictions on Test set
y_test_pred <- predict(mlr_BS, X_test_scaled)

# Model accuracy on Test set
cat("Test set accuracy: ", mean(y_test_pred == y_test))

## Test set accuracy: 0.9775
```

Multinomial Logistic Regression (FWS)

→ Train the model using Cross-Validation error to apply Forward Selection

```
# final model after FORWARD SELECTION

# create formula to fit the model
features = paste(names(X_train[, c(14, 12, 1, 13, 9)]), collapse='+')
formula = paste(c('y_train', features), collapse='~')

# Fit the Multinomial Logistic Regression model
mlr_FS <- nnet::multinom(formula, data = X_train_scaled)
```

Features selected:
battery_power, mobile_wt,
px_height, px_width, ram



```
# Make predictions on Training set
y_train_pred <- predict(mlr_FS, X_train_scaled)

# Model accuracy on Training set
cat("Training set accuracy: ", mean(y_train_pred == y_train), "\n")
```

Training set accuracy: 0.983125

```
# Make predictions on Test set
y_test_pred <- predict(mlr_FS, X_test_scaled)

# Model accuracy on Test set
cat("Test set accuracy: ", mean(y_test_pred == y_test))
```

Test set accuracy: 0.9675

Multinomial Logistic Regression (RAM only)

→ Train the model using RAM as the only feature

```
# Fit the model  
mlr_ram <- nnet::multinom(y_train ~ ram, data = X_train_scaled)
```



```
# Make predictions on Training set  
y_train_pred <- predict(mlr_ram, X_train_scaled)  
  
# Model accuracy on Training set  
cat("Training set accuracy: ", mean(y_train_pred == y_train), "\n")
```

```
## Training set accuracy: 0.749375
```

```
# Make predictions on Test set  
y_test_pred <- predict(mlr_ram, X_test_scaled)  
  
# Model accuracy on Test set  
cat("Test set accuracy: ", mean(y_test_pred == y_test))
```

```
## Test set accuracy: 0.7875
```

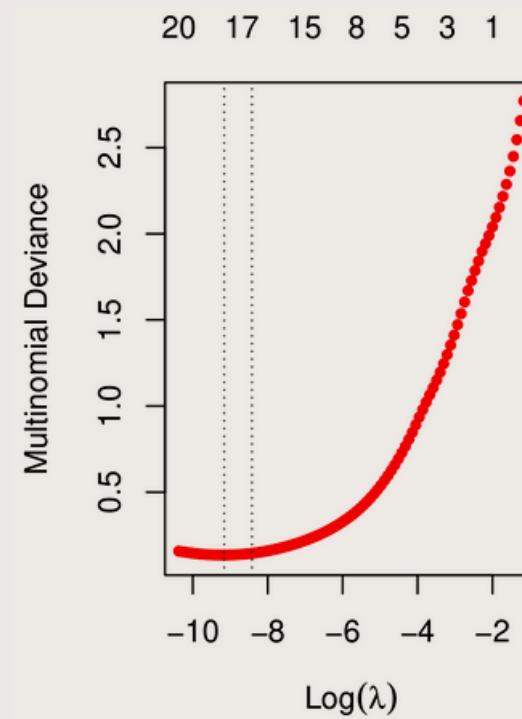
ram is not the only
influential feature

Multinomial Logistic Regression (LASSO)

→ Train the model using LASSO Regularization

```
# Lasso Regression  
mlr_lasso = cv.glmnet(as.matrix(X_train_scaled), y_train, alpha=1, family="multinomial")
```

log(λ) vs Multinomial Deviance



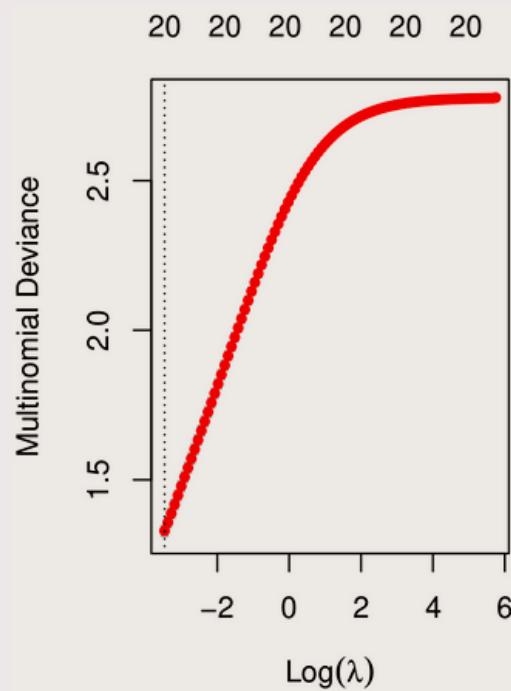
```
# LASSO predictions  
  
# Make predictions on Training set  
y_train_pred <- predict(mlr_lasso, newx = as.matrix(X_train_scaled), s = "lambda.min", type = "class")  
  
# Model accuracy on Training set  
cat("Training set accuracy with LASSO: ", mean(y_train_pred == y_train), "\n")  
  
## Training set accuracy with LASSO: 0.989375  
  
# Make predictions on Test set  
y_test_pred <- predict(mlr_lasso, newx = as.matrix(X_test_scaled), s = "lambda.min", type = "class")  
  
# Model accuracy on Test set  
cat("Test set accuracy with LASSO: ", mean(y_test_pred == y_test))  
  
## Test set accuracy with LASSO: 0.975
```

Multinomial Logistic Regression (RIDGE)

→ Train the model using RIDGE Regularization

```
# Ridge Regression  
mlr_ridge = cv.glmnet(as.matrix(X_train_scaled), y_train, alpha=0, family="multinomial")
```

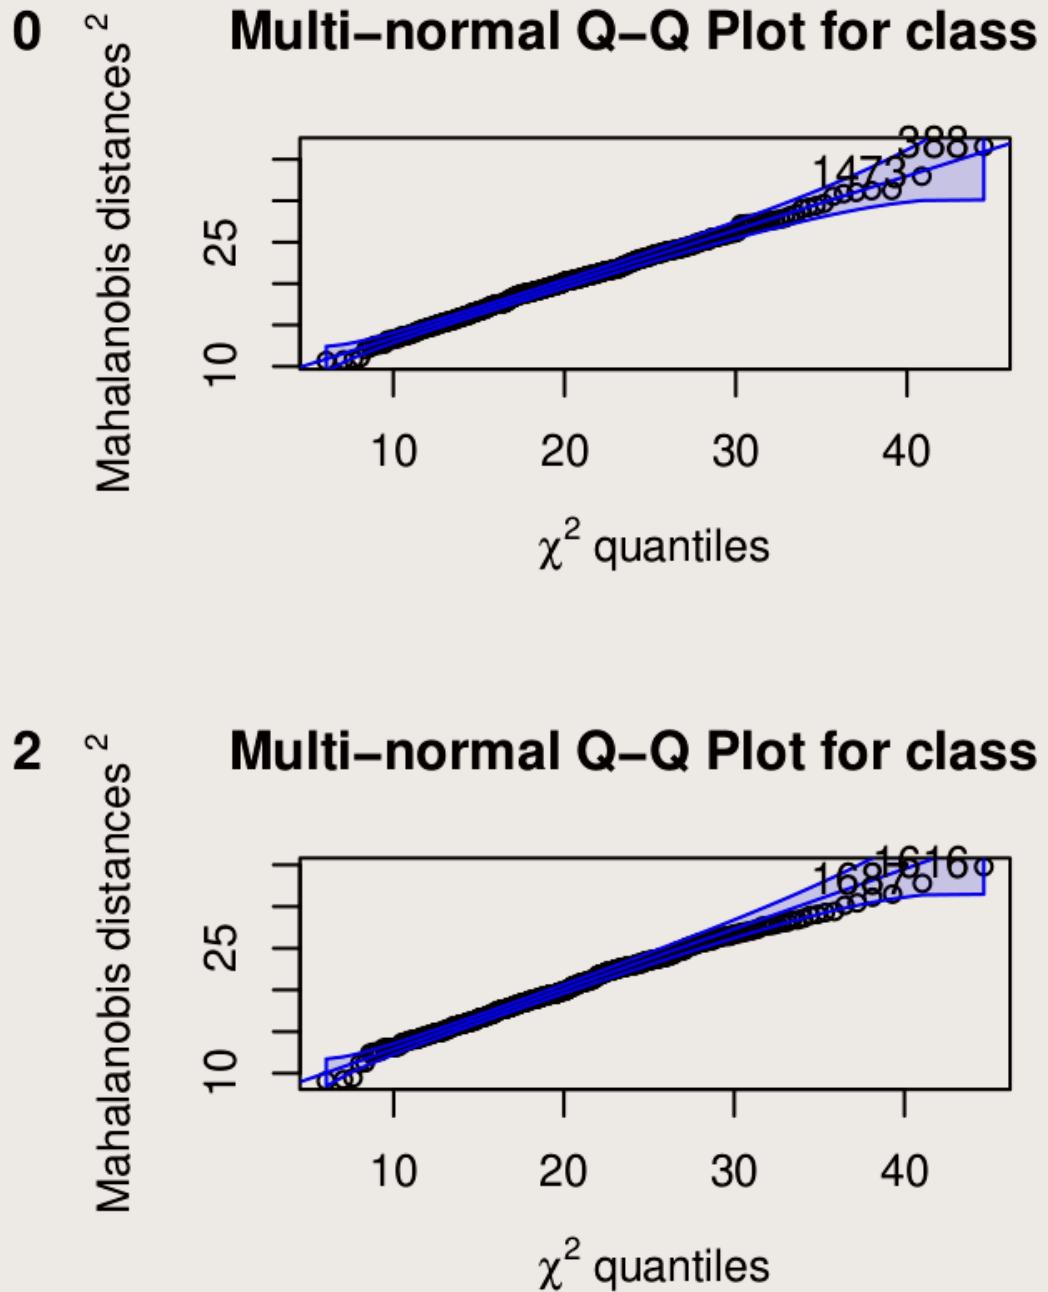
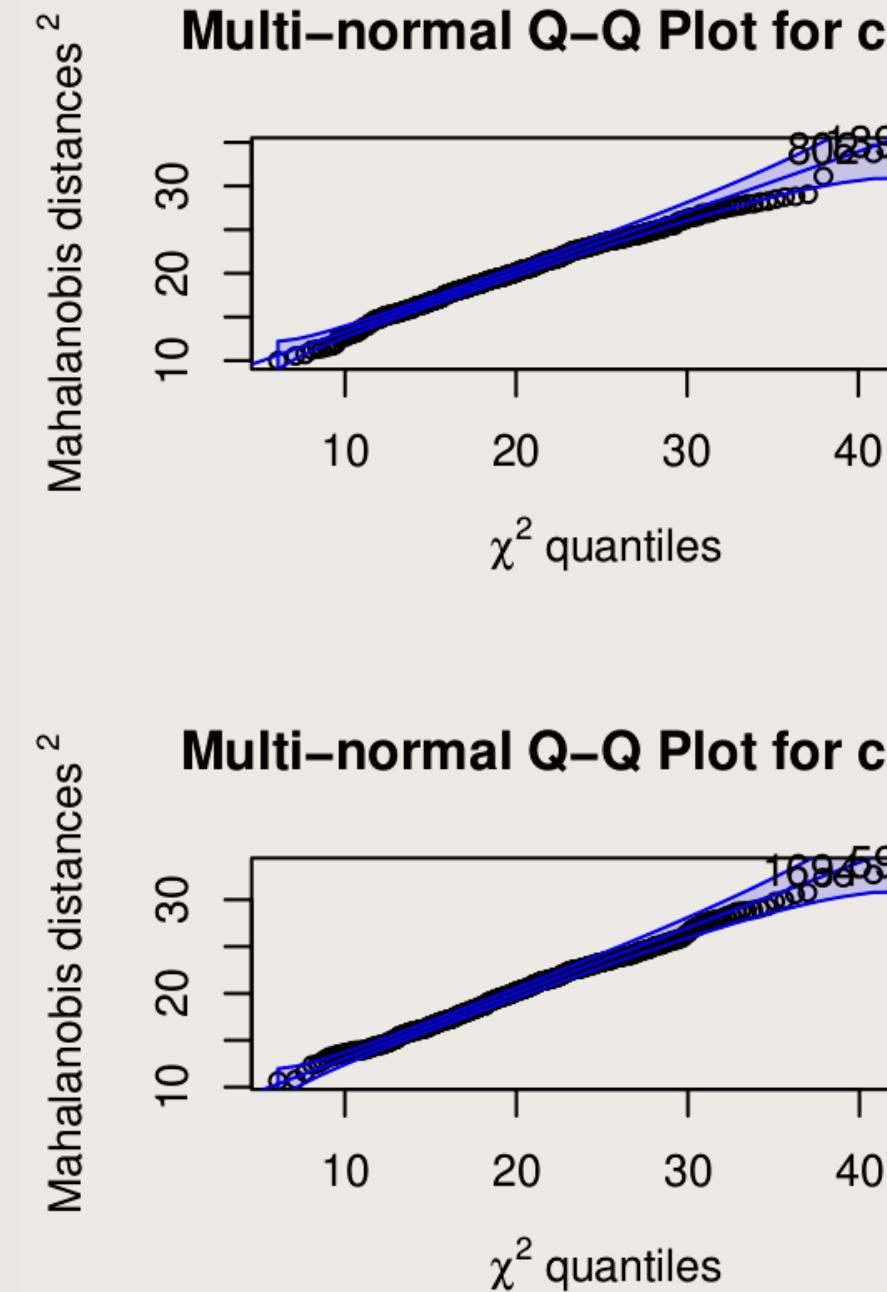
log(λ) vs Multinomial Deviance



```
# RIDGE predictions  
  
# Make predictions on Training set  
y_train_pred <- predict(mlr_ridge, newx = as.matrix(X_train_scaled), s = "lambda.min", type = "class")  
  
# Model accuracy on Training set  
cat("Training set accuracy with RIDGE: ", mean(y_train_pred == y_train), "\n")  
  
## Training set accuracy with RIDGE: 0.875625  
  
# Make predictions on Test set  
y_test_pred <- predict(mlr_ridge, newx = as.matrix(X_test_scaled), s = "lambda.min", type = "class")  
  
# Model accuracy on Test set  
cat("Test set accuracy with RIDGE: ", mean(y_test_pred == y_test))  
  
## Test set accuracy with RIDGE: 0.8525
```

Multivariate QQ-Plots

→ Check for Normal distribution within each class



Linear Discriminant Analysis (Full Dataset)

→ Train the model keeping all features

```
# original model considering all features  
LDA <- lda(y_train~, data = X_train_scaled)
```



```
CM_train <- table(y_train_pred$class, y_train)  
CM_train
```

```
##      y_train  
##      0   1   2   3  
## 0 388   9   0   0  
## 1   11 375  25   0  
## 2    0   11 361  17  
## 3    0   0   3 400
```

```
CM_test <- table(y_test_pred$class, y_test)  
CM_test
```

```
##      y_test  
##      0   1   2   3  
## 0  99   1   0   0  
## 1   2 103   6   0  
## 2    0   1 104   4  
## 3    0   0   1  79
```

```
# Make predictions on Training set  
y_train_pred <- predict(LDA, X_train_scaled)  
  
# Model accuracy on Training set  
cat("Training set accuracy: ", mean(y_train_pred$class == y_train), "\n")
```



```
## Training set accuracy: 0.9525  
  
# Make predictions on Test set  
y_test_pred <- predict(LDA, X_test_scaled)  
  
# Model accuracy on Test set  
cat("Test set accuracy: ", mean(y_test_pred$class == y_test), "\n")  
  
## Test set accuracy: 0.9625
```

Linear Discriminant Analysis (FWS Multinom.)

→ Train the model with the same features used in the Multinomial Logistic Regression (FWS)

```
# Linear Discriminant Analysis with the same features got for the Multinomial Logistic Regression model  
  
# Fit the Linear Discriminant Analysis model  
LDA <- lda(y_train~ram+px_height+battery_power+px_width+mobile_wt, data = X_train_scaled)
```

```
# Make predictions on Training set  
y_train_pred <- predict(LDA, X_train_scaled)  
  
# Model accuracy on Training set  
cat("Training set accuracy: ", mean(y_train_pred$class == y_train), "\n")  
  
## Training set accuracy: 0.95375
```

```
# Make predictions on Test set  
y_test_pred <- predict(LDA, X_test_scaled)  
  
# Model accuracy on Test set  
cat("Test set accuracy: ", mean(y_test_pred$class == y_test), '\n')
```

```
## Test set accuracy: 0.9625
```



```
##      y_train  
##      0   1   2   3  
## 0 391  7  0  0  
## 1   8 378 26  0  
## 2   0 10 359 19  
## 3   0  0  4 398
```

```
CM_train <- table(y_train_pred$class, y_train)  
CM_train
```

```
##      y_test  
##      0   1   2   3  
## 0 98  0  0  0  
## 1   3 105 7  0  
## 2   0  0 104 5  
## 3   0  0  0 78
```

```
CM_test <- table(y_test_pred$class, y_test)  
CM_test
```

Linear Discriminant Analysis (BWS Multinom.)

→ Train the model with the same features used in the Multinomial Logistic Regression (BWS)

```
# Linear Discriminant Analysis with the same features got for the Multinomial Logistic Regression model

# Fit the Linear Discriminant Analysis model
LDA <- lda(y_train~ram+battery_power+int_memory+mobile_wt+n_cores+px_height+px_width, data = X_train_sc

# Make predictions on Training set
y_train_pred <- predict(LDA, X_train_scaled)

# Model accuracy on Training set
cat("Training set accuracy: ", mean(y_train_pred$class == y_train), "\n")

## Training set accuracy: 0.95

# Make predictions on Test set
y_test_pred <- predict(LDA, X_test_scaled)

# Model accuracy on Test set
cat("Test set accuracy: ", mean(y_test_pred$class == y_test), '\n')

## Test set accuracy: 0.9625
```

```
CM_train <- table(y_train_pred$class, y_train) CM_test <- table(y_test_pred$class, y_test)
CM_train
CM_test

##      y_train          y_test
##      0   1   2   3      0   1   2   3
## 0 391  7  0  0    0 99  0  0  0
## 1   8 377 25  0    1 2 104 7  0
## 2   0 11 359 24    2 0 1 104 5
## 3   0  0  5 393    3 0  0  0 78
```



Linear Discriminant Analysis (FWS)

→ Train the model using Cross-Validation error to apply Forward Selection

```
# final model after FORWARD SELECTION

# Fit the Linear Discriminant Analysis model
LDA <- lda(y_train~ram+battery_power+px_height+px_width+mobile_wt+wifi+sc_w, data = X_train_scaled)



Features selected:  
battery_power, int_memory,  
mobile_wt, n_cores, px_height,  
px_width, ram


```

Quadratic Discriminant Analysis (Full Dataset)

→ Train the model keeping all features

```
# original model considering all features  
QDA <- qda(y_train~, data = X_train_scaled)
```

```
CM_train <- table(y_train_pred$class, y_train)  
CM_train
```

```
CM_test <- table(y_test_pred$class, y_test)  
CM_test
```

```
##      y_train  
##      0   1   2   3  
## 0 393 12  0  0  
## 1   6 377 13  0  
## 2   0   6 363  6  
## 3   0   0 13 411
```

```
##      y_test  
##      0   1   2   3  
## 0  98  3  0  0  
## 1   3 100  5  0  
## 2   0   2 100  5  
## 3   0   0   6 78
```

```
# Make predictions on Training set  
y_train_pred <- predict(QDA, X_train_scaled)  
  
# Model accuracy on Training set  
cat("Training set accuracy: ", mean(y_train_pred$class == y_train), "\n")  
  
## Training set accuracy: 0.965
```

```
# Make predictions on Test set  
y_test_pred <- predict(QDA, X_test_scaled)  
  
# Model accuracy on Test set  
cat("Test set accuracy: ", mean(y_test_pred$class == y_test), "\n")  
  
## Test set accuracy: 0.94
```



Quadratic Discriminant Analysis (FWS Multinom.)

→ Train the model with the same features used in the Multinomial Logistic Regression (FWS)

```
# Quadratic Discriminant Analysis with the same features got for the Multinomial Logistic Regression  
  
# Fit the Linear Discriminant Analysis model  
QDA <- qda(y_train~ram+px_height+battery_power+px_width+mobile_wt, data = X_train_scaled)
```

| CM_train | | CM_test | |
|----------|--------------|-----------|-------------|
| | ## y_train | ## y_test | |
| ## | 0 1 2 3 | ## | 0 1 2 3 |
| ## | 0 395 12 0 0 | ## | 0 100 2 0 0 |
| ## | 1 4 374 10 0 | ## | 1 1 101 4 0 |
| ## | 2 0 9 370 4 | ## | 2 0 2 105 4 |
| ## | 3 0 0 9 413 | ## | 3 0 0 2 79 |

```
# Make predictions on Training set  
y_train_pred <- predict(QDA, X_train_scaled)  
  
# Model accuracy on Training set  
cat("Training set accuracy: ", mean(y_train_pred$class == y_train), "\n")  
  
## Training set accuracy: 0.97  
  
# Make predictions on Test set  
y_test_pred <- predict(QDA, X_test_scaled)  
  
# Model accuracy on Test set  
cat("Test set accuracy: ", mean(y_test_pred$class == y_test), '\n')  
  
## Test set accuracy: 0.9625
```



Quadratic Discriminant Analysis (BWS Multinom.)

→ Train the model with the same features used in the Multinomial Logistic Regression (BWS)

```
# Quadratic Discriminant Analysis with the same features got for the Multinomial Logistic Regression mo  
  
# Fit the Linear Discriminant Analysis model  
QDA <- qda(y_train~ram+battery_power+int_memory+mobile_wt+n_cores+px_height+px_width, data = X_train_sc  
  
# Make predictions on Training set  
y_train_pred <- predict(QDA, X_train_scaled)  
  
# Model accuracy on Training set  
cat("Training set accuracy: ", mean(y_train_pred$class == y_train), "\n")  
## Training set accuracy: 0.965  
  
# Make predictions on Test set  
y_test_pred <- predict(QDA, X_test_scaled)  
  
# Model accuracy on Test set  
cat("Test set accuracy: ", mean(y_test_pred$class == y_test), '\n')  
## Test set accuracy: 0.9625
```

CM_train <- table(y_train_pred\$class, y_train)
CM_train

| | 0 | 1 | 2 | 3 |
|---|-----|-----|-----|-----|
| 0 | 393 | 15 | 0 | 0 |
| 1 | 6 | 373 | 11 | 0 |
| 2 | 0 | 7 | 366 | 5 |
| 3 | 0 | 0 | 12 | 412 |

CM_test <- table(y_test_pred\$class, y_test)
CM_test

| | 0 | 1 | 2 | 3 |
|---|----|-----|-----|----|
| 0 | 99 | 1 | 0 | 0 |
| 1 | 2 | 102 | 4 | 0 |
| 2 | 0 | 2 | 103 | 2 |
| 3 | 0 | 0 | 4 | 81 |



Quadratic Discriminant Analysis (FWS LDA)

→ Train the model with the same features used in the Linear Discriminant Analysis (FWS)

```
# Quadratic Discriminant Analysis with the same features got for LDA model, through Forward Selection  
  
# Fit the Linear Discriminant Analysis model  
QDA <- qda(y_train~ram+battery_power+px_height+px_width+mobile_wt+wifi+sc_w, data = X_train_scaled)
```

```
# Make predictions on Training set  
y_train_pred <- predict(QDA, X_train_scaled)  
  
# Model accuracy on Training set  
cat("Training set accuracy: ", mean(y_train_pred$class == y_train), "\n")  
  
## Training set accuracy: 0.96625  
  
# Make predictions on Test set  
y_test_pred <- predict(QDA, X_test_scaled)  
  
# Model accuracy on Test set  
cat("Test set accuracy: ", mean(y_test_pred$class == y_test), '\n')  
  
## Test set accuracy: 0.96
```



| | | CM_train | | | | CM_test | | | | | |
|----------|--|----------|---------|-----|-----|---------|--------|----|-----|-----|-----|
| | | ## | y_train | ## | ## | ## | y_test | ## | ## | | |
| | | ## | 0 | 1 | 2 | 3 | ## | 0 | 1 | 2 | 3 |
| CM_train | | ## | 0 | 395 | 17 | 0 | ## | 0 | 100 | 3 | 0 |
| | | ## | 1 | 4 | 369 | 7 | ## | 1 | 1 | 101 | 4 |
| | | ## | 2 | 0 | 9 | 370 | ## | 2 | 0 | 1 | 103 |
| | | ## | 3 | 0 | 0 | 12 | ## | 3 | 0 | 0 | 80 |
| | | ## | | | | 412 | ## | | | | |



Quadratic Discriminant Analysis (FWS)

→ Train the model using Cross-Validation error to apply Forward Selection

```
# final model after FORWARD SELECTION

# Fit the Linear Discriminant Analysis model
QDA <- qda(y_train ~ ram+battery_power+clock_speed+mobile_wt+px_height+px_width+fc, data = X_train_scaled)

# Make predictions on Training set
y_train_pred <- predict(QDA, X_train_scaled)

# Model accuracy on Training set
cat("Training set accuracy: ", mean(y_train_pred$class == y_train), "\n")

## Training set accuracy: 0.966875

# Make predictions on Test set
y_test_pred <- predict(QDA, X_test_scaled)

# Model accuracy on Test set
cat("Test set accuracy: ", mean(y_test_pred$class == y_test), '\n')

## Test set accuracy: 0.9575
```

CM_train <- table(y_train_pred\$class, y_train) CM_train

CM_test <- table(y_test_pred\$class, y_test) CM_test

| | ## | y_train | ## | y_test |
|----|-----|---------|----|--------|
| ## | 0 | 1 | ## | 0 |
| ## | 393 | 11 | ## | 100 |
| ## | 1 | 376 | ## | 1 |
| ## | 6 | 10 | ## | 1 |
| ## | 2 | 0 | ## | 0 |
| ## | 8 | 366 | ## | 2 |
| ## | 3 | 5 | ## | 105 |
| ## | 0 | 13 | ## | 0 |
| ## | 0 | 412 | ## | 2 |
| ## | 0 | | ## | 79 |

Features selected:
battery_power, clock_speed,
mobile_wt, px_height,
px_width, fc, ram

K - Nearest Neighbors (Not Scaled)

→ Train the model using original data

```
# train the K-Nearest Neighbors with the original Dataset
```

```
set.seed(1)
ctrl <- trainControl(method="repeatedcv", repeats = 5)
knn <- train(as.factor(price_range) ~ ., data = Training_set, method = "knn",
             trControl = ctrl, tuneLength = 20)
```

```
y_train_pred <- predict(knn,newdata = Training_set)

# Model accuracy on Training set
cat("Training set accuracy: ", mean(y_train_pred == as.factor(y_train)), '\n')

## Training set accuracy: 0.945625

y_test_pred <- predict(knn,newdata = Test_set)

# Model accuracy on Test set
cat("\nTest set accuracy: ", mean(y_test_pred == as.factor(y_test)), '\n')

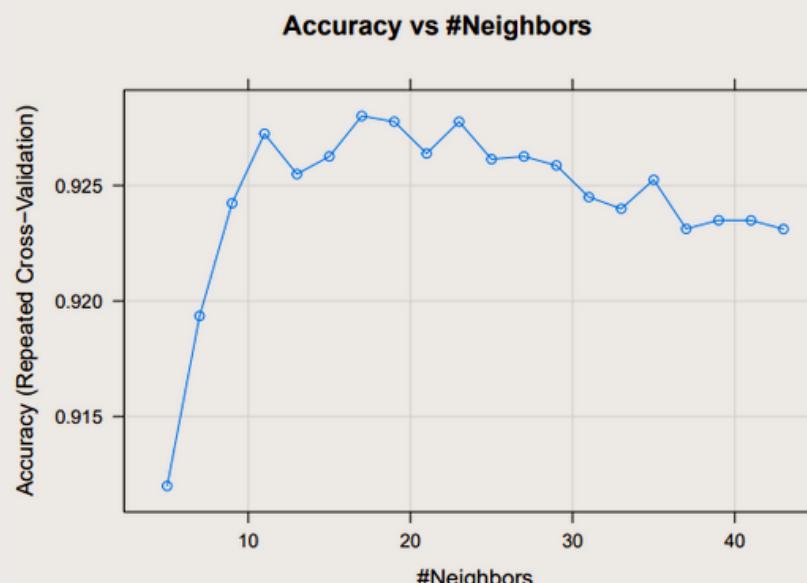
## Test set accuracy: 0.9475
```

```
table(y_train_pred, as.factor(y_train))
```

```
##
##   y_train_pred  0   1   2   3
##   0 391  18  0  0
##   1  8 367 20  0
##   2  0 10 354 16
##   3  0  0 15 401
```

```
table(y_test_pred, as.factor(y_test))
```

```
##
##   y_test_pred  0   1   2   3
##   0 98  2  0  0
##   1  3 102 4  0
##   2  0 1 101 5
##   3  0  0 6 78
```



Final value
K = 17

K - Nearest Neighbors (Scaled)

→ Train the model using scaled data

```
# train the K-Nearest Neighbors with scaled data

set.seed(1)
ctrl <- trainControl(method="repeatedcv", repeats = 5)
knn <- train(as.factor(price_range) ~ ., data = Training_set, method = "knn", trControl = ctrl,
tuneLength = 20, preProcess = c("center", "scale"))

# Model accuracy on Training set
cat("Training set accuracy: ", mean(y_train_pred == as.factor(y_train)), '\n')

## Training set accuracy: 0.684375

y_test_pred <- predict(knn, newdata = Test_set)

# Model accuracy on Test set
cat("\nTest set accuracy: ", mean(y_test_pred == as.factor(y_test)), '\n')

## Test set accuracy: 0.66675
```

table(y_train_pred, as.factor(y_train))

table(y_test_pred, as.factor(y_test))

Accuracy vs #Neighbors

Final value
K = 41

Conclusion

→ Summary of trained models

| Model | Train set accuracy (%) | Test set accuracy (%) |
|--------------------------------------------------------|-------------------------------|------------------------------|
| <i>Multinomial Logistic Regression</i> (Full Dataset) | 99.50 | 97.50 |
| <i>Multinomial Logistic Regression</i> (BWS) | 98.50 | 97.75 |
| <i>Multinomial Logistic Regression</i> (FWS) | 98.31 | 96.75 |
| <i>Multinomial Logistic Regression</i> (ram only) | 74.94 | 78.75 |
| <i>Multinomial Logistic Regression</i> (Lasso) | 98.94 | 97.50 |
| <i>Multinomial Logistic Regression</i> (Ridge) | 87.56 | 85.25 |
| <i>Linear Discriminant Analysis</i> (Full Dataset) | 95.25 | 96.25 |
| <i>Linear Discriminant Analysis</i> (FWS Multinom.) | 95.38 | 96.25 |
| <i>Linear Discriminant Analysis</i> (BWS Multinom.) | 95.00 | 96.25 |
| <i>Linear Discriminant Analysis</i> (FWS) | 95.88 | 97.50 |
| <i>Quadratic Discriminant Analysis</i> (Full Dataset) | 96.50 | 94.00 |
| <i>Quadratic Discriminant Analysis</i> (FWS Multinom.) | 97.00 | 96.25 |
| <i>Quadratic Discriminant Analysis</i> (BWS Multinom.) | 96.50 | 96.25 |
| <i>Quadratic Discriminant Analysis</i> (FWS LDA) | 96.63 | 96.00 |
| <i>Quadratic Discriminant Analysis</i> (FWS) | 96.69 | 95.75 |
| <i>K-NN</i> (not scaled) | 94.56 | 94.75 |
| <i>K-NN</i> (scaled) | 68.44 | 66.75 |

Conclusion

→ Summary of trained models

| Model | Train set accuracy (%) | Test set accuracy (%) |
|--------------------------------------------------------|------------------------|-----------------------|
| <i>Multinomial Logistic Regression</i> (Full Dataset) | 99.50 | 97.50 |
| <i>Multinomial Logistic Regression</i> (BWS) | 98.50 | 97.75 |
| <i>Multinomial Logistic Regression</i> (FWS) | 98.31 | 96.75 |
| <i>Multinomial Logistic Regression</i> (ram only) | 74.94 | 78.75 |
| <i>Multinomial Logistic Regression</i> (Lasso) | 98.94 | 97.50 |
| <i>Multinomial Logistic Regression</i> (Ridge) | 87.56 | 85.25 |
| <i>Linear Discriminant Analysis</i> (Full Dataset) | 95.25 | 96.25 |
| <i>Linear Discriminant Analysis</i> (FWS Multinom.) | 95.38 | 96.25 |
| <i>Linear Discriminant Analysis</i> (BWS Multinom.) | 95.00 | 96.25 |
| <i>Linear Discriminant Analysis</i> (FWS) | 95.88 | 97.50 |
| <i>Quadratic Discriminant Analysis</i> (Full Dataset) | 96.50 | 94.00 |
| <i>Quadratic Discriminant Analysis</i> (FWS Multinom.) | 97.00 | 96.25 |
| <i>Quadratic Discriminant Analysis</i> (BWS Multinom.) | 96.50 | 96.25 |
| <i>Quadratic Discriminant Analysis</i> (FWS LDA) | 96.63 | 96.00 |
| <i>Quadratic Discriminant Analysis</i> (FWS) | 96.69 | 95.75 |
| <i>K-NN</i> (not scaled) | 94.56 | 94.75 |
| <i>K-NN</i> (scaled) | 68.44 | 66.75 |

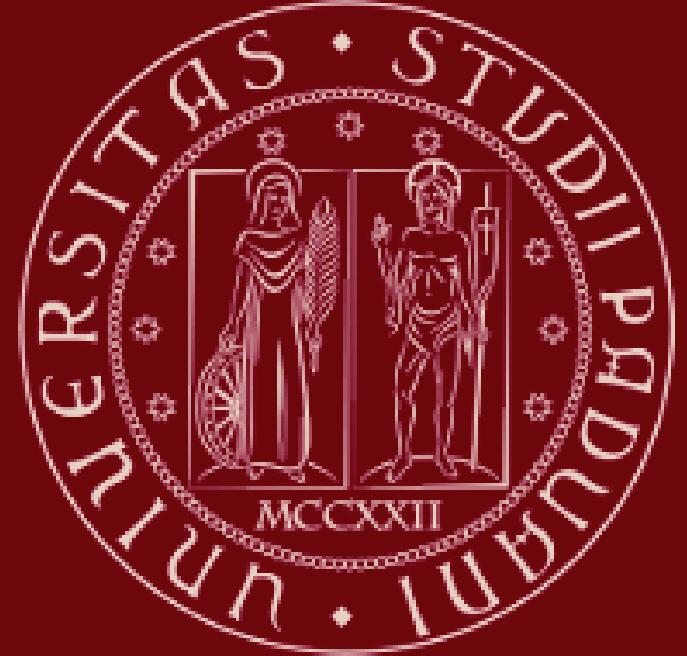
Highest
Test accuracy

Conclusion

→ Summary of trained models

| Model | Train set accuracy (%) | Test set accuracy (%) |
|-------------------------------------------------|------------------------|-----------------------|
| Multinomial Logistic Regression (Full Dataset) | 99.50 | 97.50 |
| Multinomial Logistic Regression (BWS) | 98.50 | 97.75 |
| Multinomial Logistic Regression (FWS) | 98.31 | 96.75 |
| Multinomial Logistic Regression (ram only) | 74.94 | 78.75 |
| Multinomial Logistic Regression (Lasso) | 98.94 | 97.50 |
| Multinomial Logistic Regression (Ridge) | 87.56 | 85.25 |
| Linear Discriminant Analysis (Full Dataset) | 95.25 | 96.25 |
| Linear Discriminant Analysis (FWS Multinom.) | 95.38 | 96.25 |
| Linear Discriminant Analysis (BWS Multinom.) | 95.00 | 96.25 |
| Linear Discriminant Analysis (FWS) | 95.88 | 97.50 |
| Quadratic Discriminant Analysis (Full Dataset) | 96.50 | 94.00 |
| Quadratic Discriminant Analysis (FWS Multinom.) | 97.00 | 96.25 |
| Quadratic Discriminant Analysis (BWS Multinom.) | 96.50 | 96.25 |
| Quadratic Discriminant Analysis (FWS LDA) | 96.63 | 96.00 |
| Quadratic Discriminant Analysis (FWS) | 96.69 | 95.75 |
| K-NN (not scaled) | 94.56 | 94.75 |
| K-NN (scaled) | 68.44 | 66.75 |

Other
relevant values



Thanks for the attention!